

COMP1002 Final Assignment Report

BY THEJANA KOTTAWATTA HEWAGE
22307822

1) Overview: System Purpose and Modular Design

The purpose of the project is to optimise hospital workflows. Each module of the project is constructed in a way which offer best performance for each module's task.

1. For module 1, a graph-based adjacency list is chosen as the ideal data structure for modelling networks. It represents hospital's each department as vertices and corridors as edges.
2. For module 2, a has table is implemented to provide an average case time complexity for inserting and searching of patient records.
3. For module 3, a max heap is used to assign priority to patients. This module of the program ensures that highest priority patient is identified and treated more efficiently.
4. For module 4, merge sort and quick sort are used to sort large datasets.

2) Data Structures Used: Justification for graph, hash table, heap, sorting.

Module 2: Hash-Based Patient Lookup – Short written justification of hash strategy, table size and hash function

1. Collision Handling Strategy: Chaining

The collision handling strategy I've chosen is chaining using linked lists, over open addressing methods like linear probing.

- Chaining is generally simpler to implement, particularly when it comes to the deletion of entries. In open addressing DELETE often requires complex logic to maintain the integrity of subsequent lookups. Whereas in chaining, you simply remove a node from the linked list.
- Chaining handles high load factors (many entries compared to the table size) more gracefully. In contrast, open addressing methods suffer from clustering, where occupied slots form large contiguous blocks, causing performance to degrade across the entire table.

2. Hash Function & Table Size strategy: Modulo + Prime Number

- The primary goal of the hash function I have chosen is to minimise collisions and apply uniform key distribution.
 - Hash function: The sum of ASCII values of the patient ID's is used for has I believe this was sufficient for the application's needs.
 - Table Size optimisation: The table size is always set to the next prime number greater than the requested capacity.
 - Prime modulo operator: Remainder (%) of when hash is divided by size of the table is used to
- Why Prime? Using a prime number as the table size is crucial when the modulo operator (the remainder of the hash divided by the table size) is used to map the hash value to a bucket index. A prime size helps ensure that the resulting indices are more uniformly distributed, which significantly reduces the number of initial collisions.

Module 3: Heap-Based Emergency Scheduling – Brief written note on update strategy

Update Strategy: Reinsertion is the strategy I have chosen for handling a change in a patient's urgency level. Each update is treated as a new request. A potential new entry with the new, re-calculated priority is inserted into the heap. My approach avoids any alternative approaches such as: a manual search & update O(n) operation. When the old, lower-priority request is eventually found, the system can perform a quick check against the hash table to see if the patient's status has changed or if they have already been treated and discard the old request.

Implementation: 0-Based Array Per the provided example, this heap uses a 0-based array.

Parent of i: $(i - 1) // 2$

Left Child of i: $2 * i + 1$

Right Child of i: $2 * i + 2$

Module 4: Sorting Patient Records - Brief write-up on stability and pivot strategy.

Merge Sort: Top-down recursive approach is an intuitive way. Moreover, it uses a temporary array to merge results back into the original array “A”.

Quicksort: Quicksort uses *median of 3* pivot strategy.

- Pivot strategy (implementation):
 - GetMedianPivot function inspects the 1st, middle and final elements.
 - Manually sort the 3 elements in and return the median value as an index.
- My chosen pivot strategy is a good against the worst-case scenario: causing an already sorted list to be resorted.

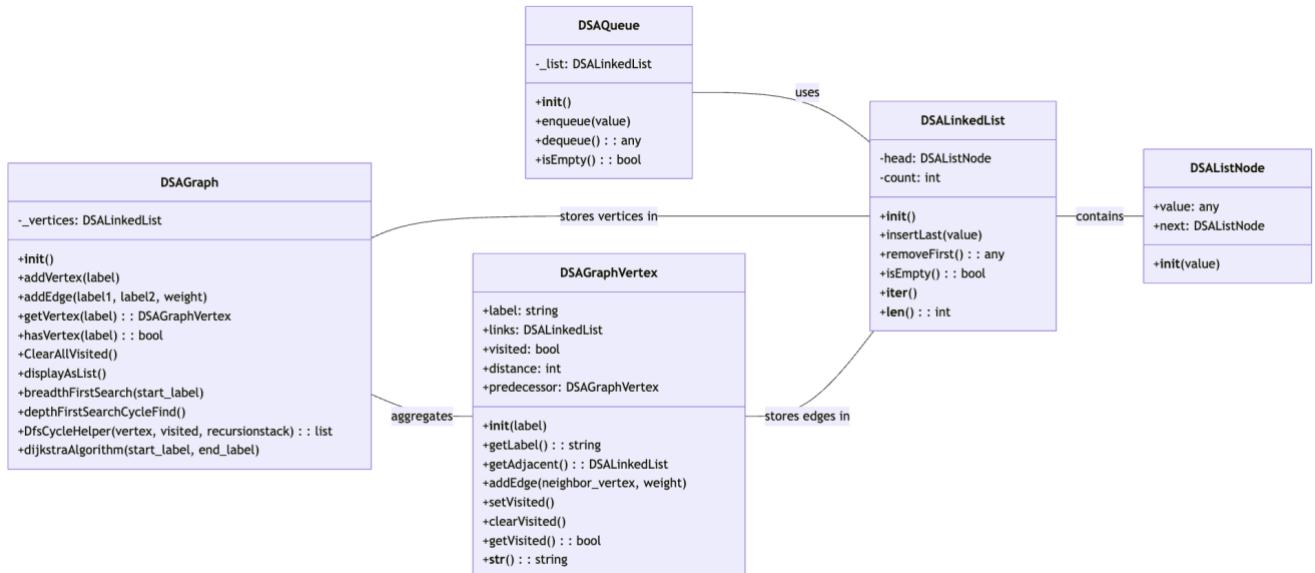
Stability:

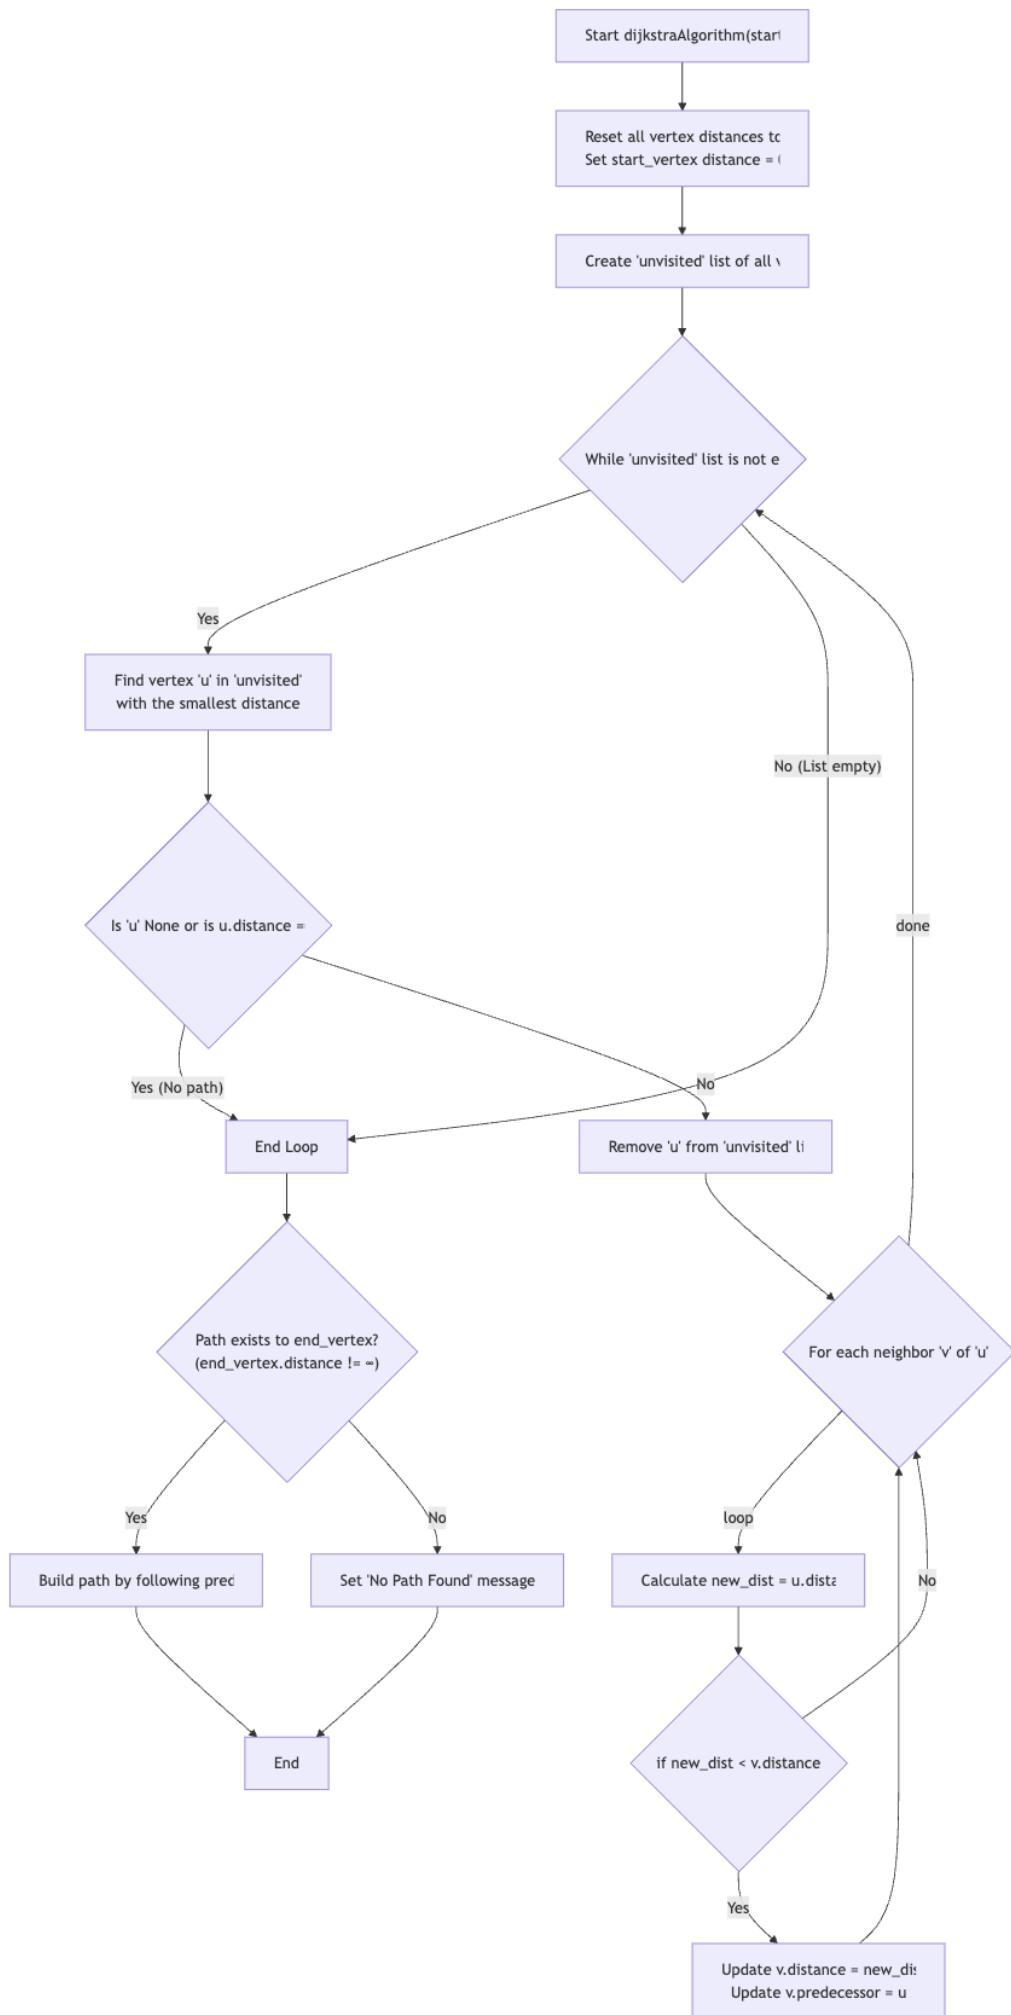
- MergeSort is very stable. During merge, if 2 elements have the same duration, element from left array is taken first, preserving the original order of equal elements.
- Quicksort is not stable. Partition involves long distance swaps which don't preserve order of elements.

3) Module Integration: Flowchart/UML describing data flow and interactions.

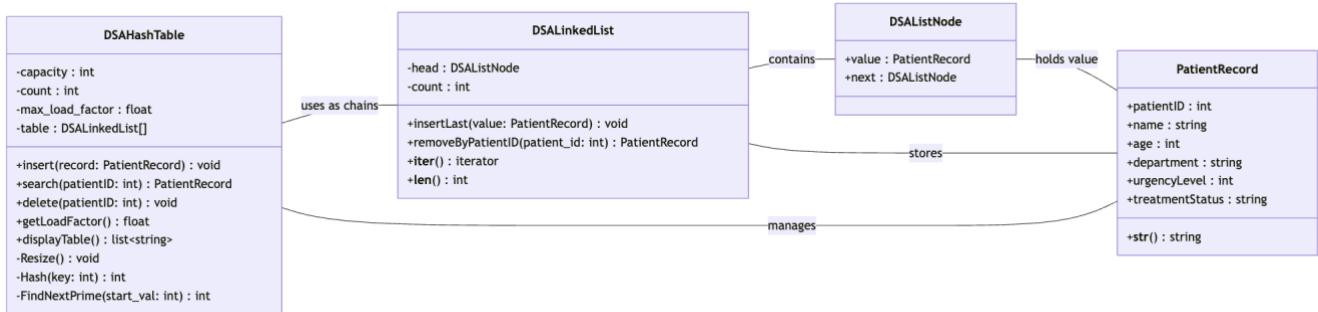
Flowchart/UML describing data flow and interactions

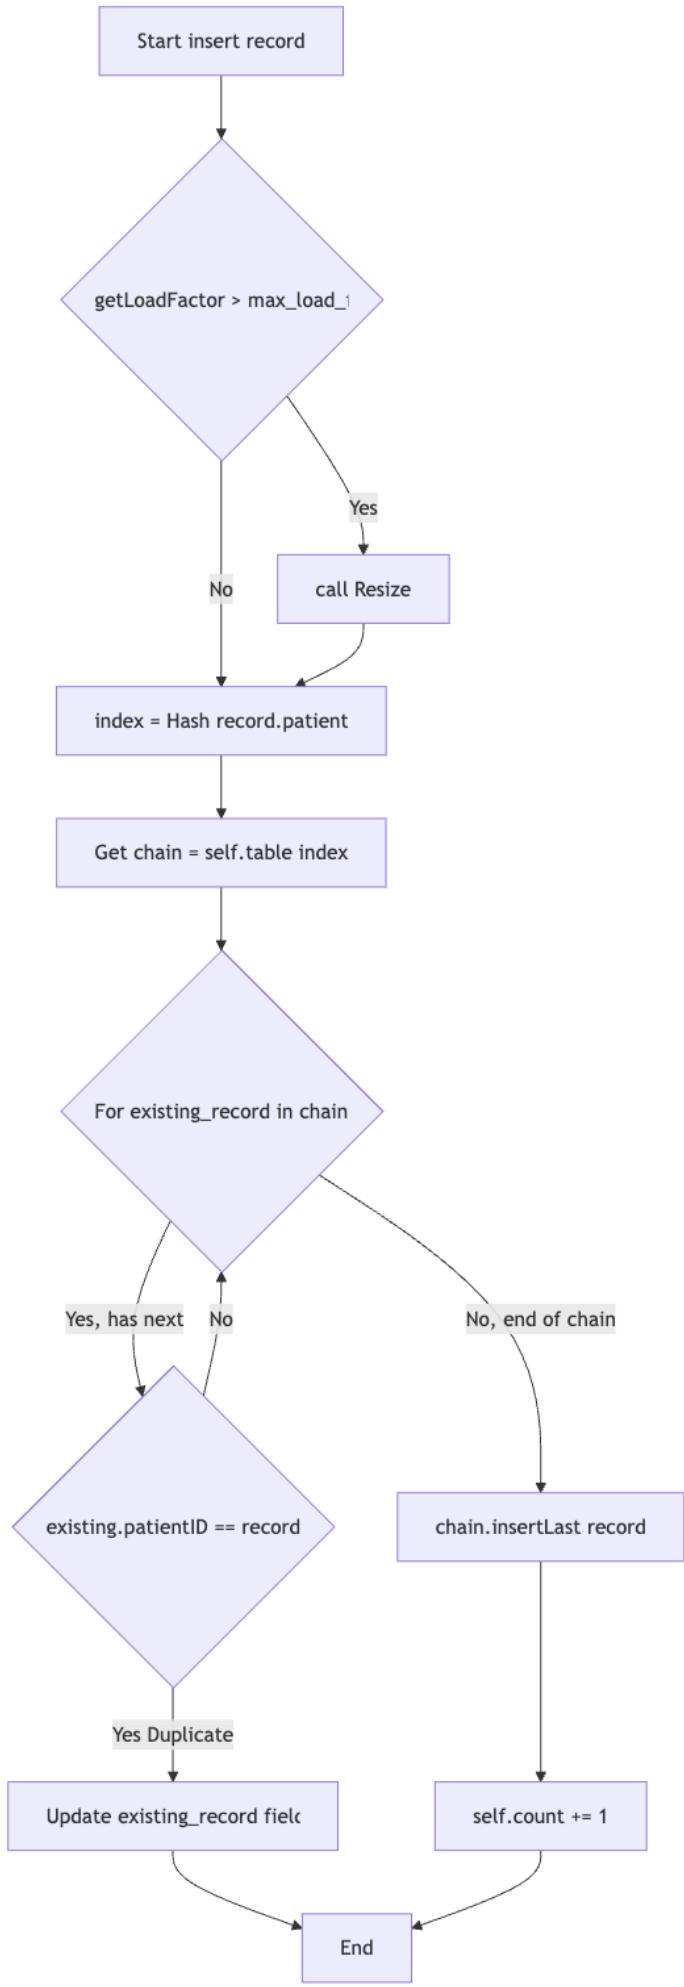
Module 1: Graph Based Hospital Navigation



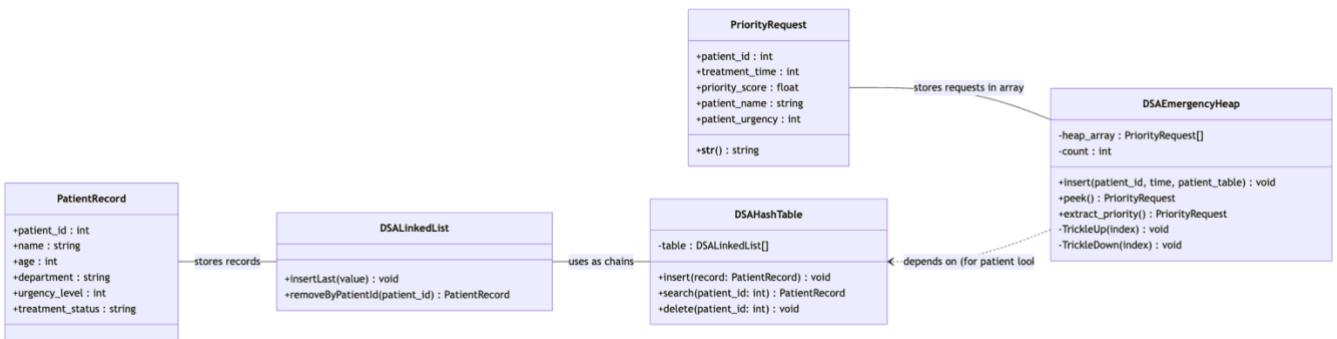


Module 2: Hash-based patient lookup

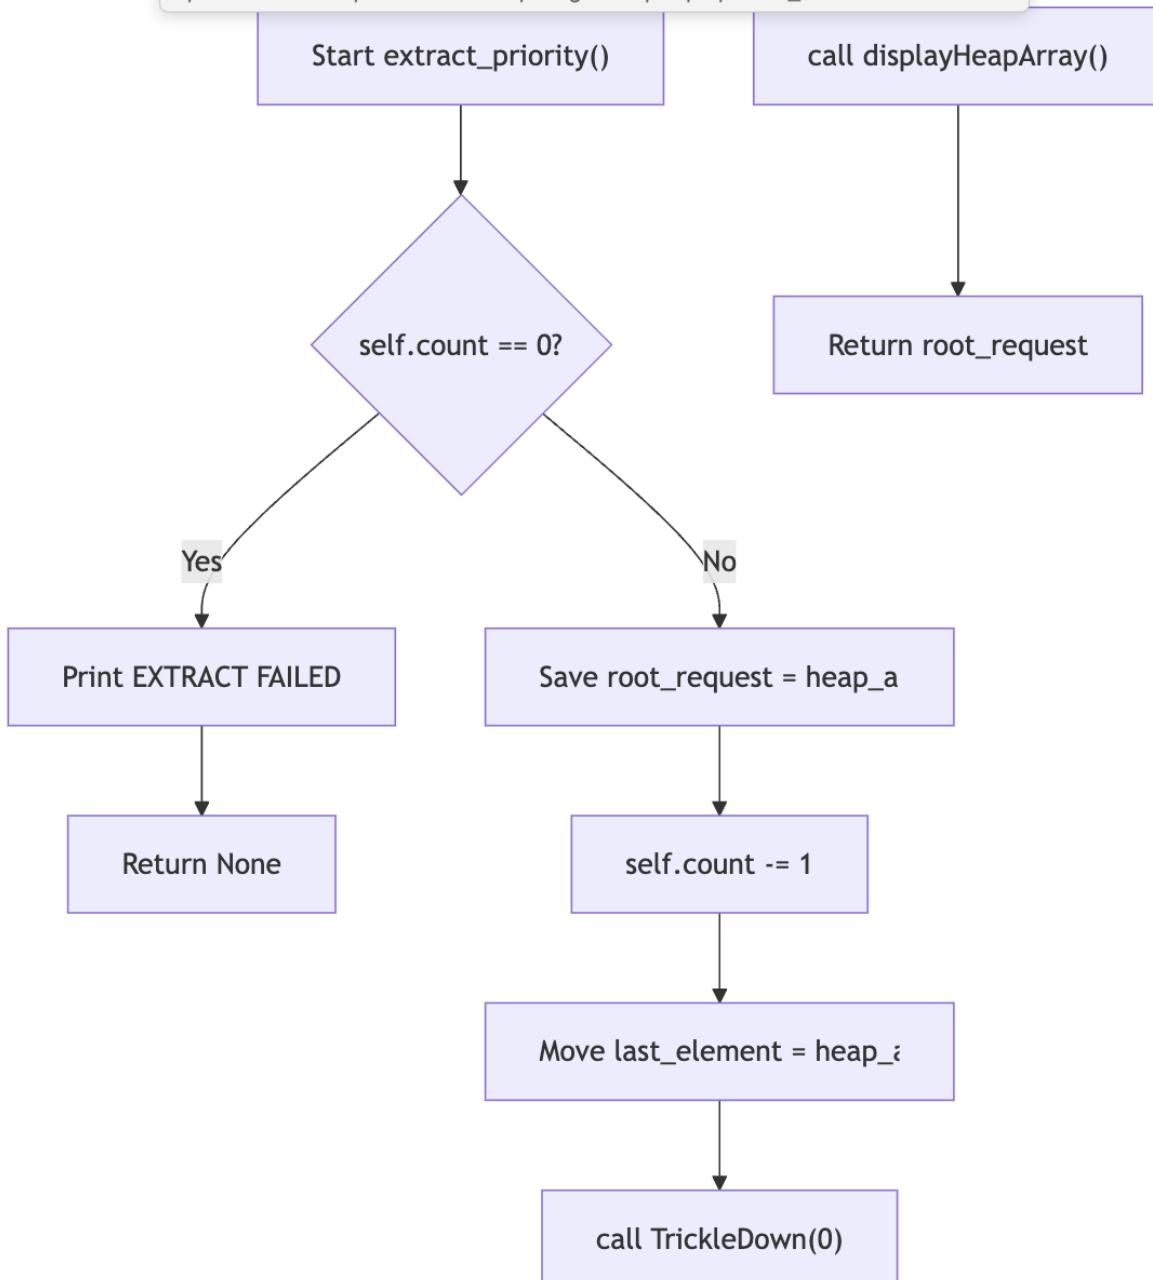




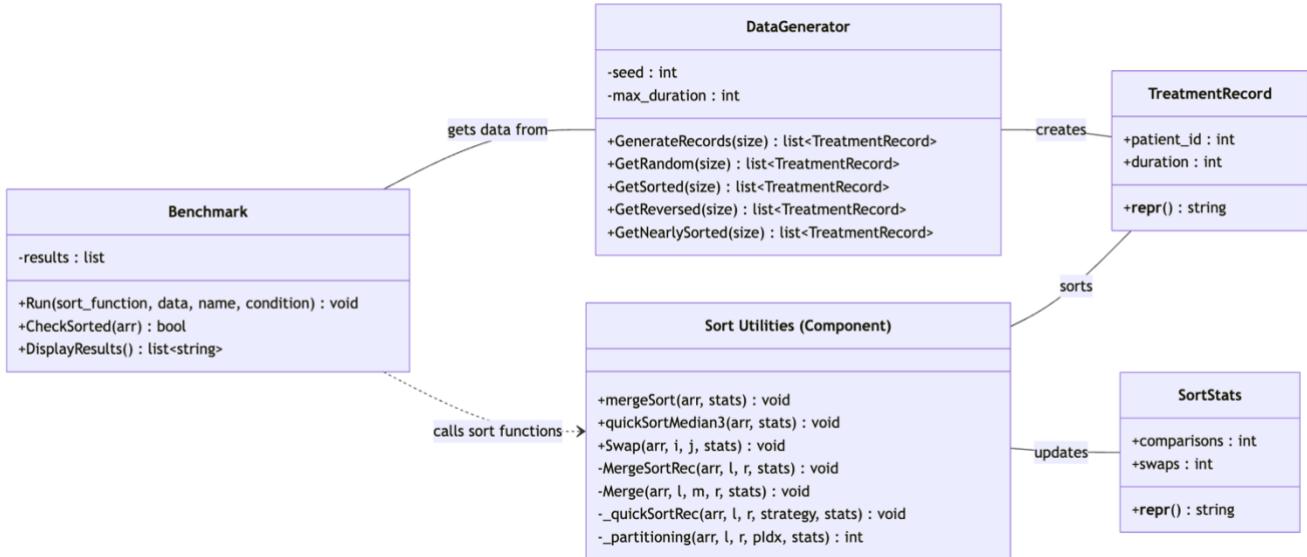
Module 3: Heap-based emergency scheduling.



~/COMP1000DSA/COMP1002DSA/Assignment/output/2hash_results.txt • Modified



Module 4: Sorting Patients Records



4) Algorithm Complexity: Time and space analysis for each module.

Module 1: Graph Based Hospital Navigation

- Performance of module 1 is defined mainly by its 2 main variables V (no. Of departments/vertices) and E (edges/corridors).
- The BFS DFS algorithms have linear time complexity.
- Dijkstra's algorithm has a time complexity of $O(n^2)$. This is a trade off because, as we are unable to use built-in functions, a manual scan of the list of departments is required.

Module 2: Hash-based patient lookup

- Module achieves an average-case time complexity $O(1)$ for `insert()`, `search()` and `delete()`.
- This is achieved with the combination of prime number and good hash function for the distribution of keys.
- The good performance is depicted by the load factor being a healthy figure of 0.45, even after undergoing resizing.
- Primary trade-off: choosing chaining for collision handling.

Module 3: Heap-based emergency scheduling.

- Core operations `insert()` and `extract_priority()` achieve a time complexity of $O(n \log n)$.
 - They both use trickle up/down and both need to traverse the height of the heap tree rather than the entire array.
- `Peek()` finds the highest priority patient with an instantaneous $O(1)$.
- Trade off: a small cost of $O(\log n)$ on every insertion.

Module 4: Sorting Patients Records

- As n increases from 100 \rightarrow 1000, runtime increases from 0.08ms to 1.3ms. ($n \log n$ growth).

- Merge sort is independent of order of input. Time on reversed data is the same as random input (1.106ms == 1.319ms).
- Conversely, quick sort operates using $O(\log n)$. With $O(n^2)$ on worst case input.

5) Sample Output: Screenshots/log excerpts per module.

Module 1: Graph Based Hospital Navigation

```

~/COMP1000DSA/COMP1002DSA/Assignment/module1.py · Deleted
1 ######
2 ### Hospital Navigation ###
3 #####
4
5
6 --- Hospital Adjacency List ---
7 Emergency      | ICU(4) -> Radiology(2) -> Wards(7)
8 ICU            | Emergency(4) -> Operating Theatres(2) -> Pharmacy(3)
9 Pharmacy        | ICU(3) -> Laboratories(2)
10 Radiology     | Emergency(2) -> Laboratories(3)
11 Laboratories   | Radiology(3) -> Pharmacy(2)
12 Operating Theatres | ICU(2) -> Wards(2)
13 Wards          | Emergency(7) -> Operating Theatres(2) -> Outpatient Units(5) -> Cafeteria(3)
14 Outpatient Units | Wards(5) -> Cafeteria(2)
15 Cafeteria      | Wards(3) -> Outpatient Units(2)
16 Morgue         |
17
18 --- BFS starting from 'Emergency' ---
19 Level 0: Emergency, ICU, Radiology, Wards
20 Level 1: Operating Theatres, Pharmacy, Laboratories, Outpatient Units, Cafeteria
21 Level 2:
22 Cycle Detected: Emergency -> ICU -> Emergency
23
24
25 --- Shortest Path from 'Emergency' to 'Outpatient Units' ---
26 Path: Emergency -> Wards -> Outpatient Units
27 Total walking time: 12 minutes.
28
29
30
31 --- Shortest Path from 'Emergency' to 'Morgue' ---
32 No path found from 'Emergency' to 'Morgue'.
33

```

Module 2: Hash-based patient lookup

```
Assignment > output > 2hash_results.txt
1 ##### MODULE 2: Hash-Based Patient Lookup #####
2 #####
3 #####
4
5 ### Step 1: Demonstrating Inserts ###
6 Demonstration: Hash(112) = 5, Hash(213) = 10
7 Patient records system is populated.
8
9
10 COLLISION DEMO: Patient 112 and 213 both hash to the same initial index.
11 | Hash Table will resolve collisions via the collision strategy I've implemented; Chaining.
12
13 ===== HASH TABLE =====
14 Count: 21, Capacity: 47, Load Factor: 0.45
15 Index 00: [ID: 745] -> None
16 Index 01: [ID: 1734] -> None
17 Index 06: [ID: 2088] -> None
18 Index 08: [ID: 901] -> None
19 Index 09: [ID: 1599] -> None
20 Index 11: [ID: 451] -> None
21 Index 12: [ID: 1366] -> None
22 Index 13: [ID: 112] -> None
23 Index 22: [ID: 1245] -> None
24 Index 23: [ID: 1010] -> None
25 Index 26: [ID: 819] -> [ID: 304] -> [ID: 1602] -> None
26 Index 28: [ID: 101] -> None
27 Index 29: [ID: 1977] -> None
28 Index 31: [ID: 1123] -> None
29 Index 34: [ID: 522] -> [ID: 1851] -> None
30 Index 35: [ID: 213] -> None
31 Index 40: [ID: 633] -> None
32 Index 44: [ID: 1482] -> None
33 =====
34
35
36 ### Step 2: Searching for Patients ###
37 Searching for an existing patient: (patient 101)
38 | Record found: ID: 101, Name: John Smith, Age: 45, Dept: Cardiology, Urgency: 2, Status: Admitted
39
40 Searching for a non-existent patient: (patient 999)
41 | Record not found (SUCCESSFUL).
42
43 ### Step 3: DELETE Demonstration ###
44 Deleting Patient 451...
45 Attempting to search for the deleted patient:
46 | Record not found (SUCCESSFULLY DELETED.)
47
48 Attempting to delete a non-existent patient:
49 | Delete attempt finished (GRACEFUL ERROR).
50
51 ===== HASH TABLE =====
52 Count: 20, Capacity: 47, Load Factor: 0.43
53 Index 00: [ID: 745] -> None
54 Index 01: [ID: 1734] -> None
55 Index 06: [ID: 2088] -> None
56 Index 08: [ID: 901] -> None
57 Index 09: [ID: 1599] -> None
58 Index 12: [ID: 1366] -> None
59 Index 13: [ID: 112] -> None
```

```
59 Index 13: [ID: 112] -> None
60 Index 22: [ID: 1245] -> None
61 Index 23: [ID: 1010] -> None
62 Index 26: [ID: 819] -> [ID: 304] -> [ID: 1602] -> None
63 Index 28: [ID: 101] -> None
64 Index 29: [ID: 1977] -> None
65 Index 31: [ID: 1123] -> None
66 Index 34: [ID: 522] -> [ID: 1851] -> None
67 Index 35: [ID: 213] -> None
68 Index 40: [ID: 633] -> None
69 Index 44: [ID: 1482] -> None
70 =====
71
72
73 --- Final Summary ---
74 Complexity: All operations (insert, search, delete) demonstrated low 'hop' counts,
75 | | | | supporting the expected O(1) average time complexity. The worst-case for a single
76 | | | | operation would be O(n) if all keys hashed to the same index.
77
78 Load Factor Behaviour: The table began with a capacity of 11 and was resized
79 | | | | to 23 when the load factor exceeded the threshold of 0.7,
80 | | | | ensuring that chains remain short and performance remains high.
```

Module 3: Heap-based emergency scheduling.

```

Assignment > output > E: 3heap_results.txt
1 ##########
2 ## (3) Emergency Scheduler ##
3 #####
4
5 ### Step 1: Initialising hash table to store Patient Records from input/patients.csv ###
6
7 RESIZING: Load factor == 0.7.
8 BEFORE: 23 -> AFTER: 47.
9
10 Patient records system is populated.
11
12 ### Step 2: Inserting 10 Emergency Requests from input/requests.csv ###
13 SEARCH SUCCESS!: Found Patient 304 at index 26.
14 INSERTING: Mary Williams (ID: 304) with U=1, T=20m. Priority = (6-1) + (1000/20) = 55.00
15 | Heap State: {[Index 0: ID 304, Priority 55.00]}
16 -----
17 SEARCH SUCCESS!: Found Patient 522 at index 34.
18 INSERTING: Susan Miller (ID: 522) with U=5, T=120m. Priority = (6-5) + (1000/120) = 9.33
19 | Heap State: {[Index 0: ID 304, Priority 55.00], [Index 1: ID 522, Priority 9.33]}
20 -----
21 SEARCH SUCCESS!: Found Patient 213 at index 35.
22 INSERTING: Jane Doe (ID: 213) with U=1, T=60m. Priority = (6-1) + (1000/60) = 21.67
23 | Heap State: {[Index 0: ID 304, Priority 55.00], [Index 1: ID 522, Priority 9.33], [Index 2: ID 213, Priority 21.67]}
24 -----
25 SEARCH SUCCESS!: Found Patient 112 at index 13.
26 INSERTING: Peter Jones (ID: 112) with U=3, T=10m. Priority = (6-3) + (1000/10) = 103.00
27 | Heap State: {[Index 0: ID 112, Priority 103.00], [Index 1: ID 304, Priority 55.00], [Index 2: ID 213, Priority 21.67], [Index 3: ID 522, Priority 9.33]}
28 -----
29 SEARCH SUCCESS!: Found Patient 101 at index 28.
30 INSERTING: John Smith (ID: 101) with U=2, T=30m. Priority = (6-2) + (1000/30) = 37.33
31 | Heap State: {[Index 0: ID 112, Priority 103.00], [Index 1: ID 304, Priority 55.00], [Index 2: ID 213, Priority 21.67], [Index 3: ID 522, Priority 9.33], [Index 4: ID 101, Priority 37.33]}
32 -----
33 SEARCH SUCCESS!: Found Patient 451 at index 11.
34 INSERTING: David Brown (ID: 451) with U=4, T=45m. Priority = (6-4) + (1000/45) = 24.22
35 | Heap State: {[Index 0: ID 112, Priority 103.00], [Index 1: ID 304, Priority 55.00], [Index 2: ID 451, Priority 24.22], [Index 3: ID 522, Priority 9.33], [Index 4: ID 101, Priority 37.33], [Index 5: ID 213, Priority 21.67]}
36 -----
37 SEARCH SUCCESS!: Found Patient 745 at index 0.
38 INSERTING: Linda Garcia (ID: 745) with U=3, T=15m. Priority = (6-1) + (1000/15) = 71.67
39 | Heap State: {[Index 0: ID 112, Priority 103.00], [Index 1: ID 304, Priority 55.00], [Index 2: ID 745, Priority 71.67], [Index 3: ID 522, Priority 9.33], [Index 4: ID 101, Priority 37.33], [Index 5: ID 213, Priority 21.67]}
40 -----
41 SEARCH SUCCESS!: Found Patient 633 at index 40.
42 INSERTING: Robert Wilson (ID: 633) with U=2, T=25m. Priority = (6-2) + (1000/25) = 44.00
43 | Heap State: {[Index 0: ID 112, Priority 103.00], [Index 1: ID 304, Priority 55.00], [Index 2: ID 745, Priority 71.67], [Index 3: ID 633, Priority 44.00], [Index 4: ID 101, Priority 37.33], [Index 5: ID 213, Priority 21.67]}
44 -----
45 SEARCH SUCCESS!: Found Patient 1602 at index 26.
46 INSERTING: Daniel Allen (ID: 1602) with U=4, T=60m. Priority = (6-4) + (1000/60) = 18.67
47 | Heap State: {[Index 0: ID 112, Priority 103.00], [Index 1: ID 304, Priority 55.00], [Index 2: ID 745, Priority 71.67], [Index 3: ID 633, Priority 44.00], [Index 4: ID 101, Priority 37.33], [Index 5: ID 213, Priority 21.67]}
48 -----
49 SEARCH SUCCESS!: Found Patient 819 at index 26.
50 INSERTING: Michael Martinez (ID: 819) with U=3, T=90m. Priority = (6-3) + (1000/90) = 14.11
51 | Heap State: {[Index 0: ID 112, Priority 103.00], [Index 1: ID 304, Priority 55.00], [Index 2: ID 745, Priority 71.67], [Index 3: ID 633, Priority 44.00], [Index 4: ID 101, Priority 37.33], [Index 5: ID 213, Priority 21.67]}
52 -----
53
54 ### Step 3: Edge Case ###
55 SEARCH FAIL: Patient 999 not found.
56 INSERT FAILED: Patient 999 not found in records.
57 SEARCH SUCCESS!: Found Patient 101 at index 28.
58 INSERT FAILED: Invalid treatment time (0) for Patient 101.
59 -----
```

```

Assignment > output > E: 3heap_results.txt
59 -----
60 (Demo: Deleting patient No. 1010 from records...)
61 DELETE: Successfully removed Patient 1010.
62 SEARCH FAIL: Patient 1010 not found.
63 INSERT FAILED: Patient 1010 not found in records.
64 -----
65
66 ### Step 4: Extracting Top 5 Priority Patients ---
67 Expected order: 112(103.0), 745(71.7), 304(55.0), 633(44.0), 101(37.3)
68 -----
69
70 Extracting patient 1...
71 EXTRACTED: [ID: 112 (Peter Jones), Priority: 103.00, U: 3, T: 10m]
72 | Heap State: {[Index 0: ID 745, Priority 71.67], [Index 1: ID 304, Priority 55.00], [Index 2: ID 451, Priority 24.22], [Index 3: ID 633, Priority 44.00], [Index 4: ID 101, Priority 37.33], [Index 5: ID 213, Priority 21.67]}
73 -----
74
75 Extracting patient 2...
76 EXTRACTED: [ID: 745 (Linda Garcia), Priority: 71.67, U: 1, T: 15m]
77 | Heap State: {[Index 0: ID 304, Priority 55.00], [Index 1: ID 451, Priority 24.22], [Index 2: ID 1602, Priority 18.67], [Index 3: ID 101, Priority 37.33], [Index 4: ID 213, Priority 21.67]}
78 -----
79
80 Extracting patient 3...
81 EXTRACTED: [ID: 304 (Mary Williams), Priority: 55.00, U: 1, T: 20m]
82 | Heap State: {[Index 0: ID 633, Priority 44.00], [Index 1: ID 101, Priority 37.33], [Index 2: ID 451, Priority 24.22], [Index 3: ID 1602, Priority 18.67], [Index 4: ID 522, Priority 9.33], [Index 5: ID 213, Priority 21.67]}
83 -----
84
85 Extracting patient 4...
86 EXTRACTED: [ID: 633 (Robert Wilson), Priority: 44.00, U: 2, T: 25m]
87 | Heap State: {[Index 0: ID 101, Priority 37.33], [Index 1: ID 1602, Priority 18.67], [Index 2: ID 451, Priority 24.22], [Index 3: ID 819, Priority 14.11], [Index 4: ID 522, Priority 9.33], [Index 5: ID 213, Priority 21.67]}
88 -----
89
90 Extracting patient 5...
91 EXTRACTED: [ID: 101 (John Smith), Priority: 37.33, U: 2, T: 30m]
92 | Heap State: {[Index 0: ID 451, Priority 24.22], [Index 1: ID 1602, Priority 18.67], [Index 2: ID 213, Priority 21.67], [Index 3: ID 819, Priority 14.11], [Index 4: ID 522, Priority 9.33]}
93 -----
94
95 ### Final Summary ###
96 Evidence: The extraction log clearly shows that patients with the highest computed
97 priority scores (e.g., 103.0, 71.7) were extracted first, regardless of the order they
98 were inserted in. This confirms the Max Heap function is serving higher-priority patients first consistently.
99 -----
```

Module 4: Sorting Patients Records

Assignment > output > 4benchmark_results.txt

```
1 #####  
2 ### MODULE 4: Sorting Patients Records ###  
3 #####  
4  
5 --- 1. Sample Sort Output (Size=100, Random) ---  
6 Original (first 10): ['205m', '180m', '128m', '95m', '207m', '155m', '293m', '151m', '206m', '14m']  
7 Sorted (first 10): ['14m', '17m', '19m', '19m', '21m', '21m', '23m', '24m', '26m', '30m']  
8 Sorted (last 10): ['247m', '253m', '262m', '280m', '282m', '286m', '287m', '292m', '293m', '294m']  
9  
10 --- 2. Running Full Benchmark ---  
11 Generating and testing datasets for size 100...  
12 Generating and testing datasets for size 500...  
13 Generating and testing datasets for size 1000...  
14  
15 ======  
16 --- Sorting Algorithm Benchmark Results ---  
17 ======  
18 Algorithm | Condition | Size | Time (ms) | Comparisons | Swaps | Passed  
19 -----  
20 Merge Sort | Random | 100 | 0.100 | 542 | 0 | Yes  
21 Quick Sort | Random | 100 | 0.102 | 707 | 482 | Yes  
22 Merge Sort | Reversed | 100 | 0.085 | 322 | 0 | Yes  
23 Quick Sort | Reversed | 100 | 0.103 | 839 | 531 | Yes  
24 Merge Sort | Nearly Sorted | 100 | 0.095 | 502 | 0 | Yes  
25 Quick Sort | Nearly Sorted | 100 | 0.091 | 798 | 419 | Yes  
26 Merge Sort | Random | 500 | 0.656 | 3852 | 0 | Yes  
27 Quick Sort | Random | 500 | 0.830 | 5446 | 3233 | Yes  
28 Merge Sort | Reversed | 500 | 0.558 | 2371 | 0 | Yes  
29 Quick Sort | Reversed | 500 | 0.872 | 6702 | 4162 | Yes  
30 Merge Sort | Nearly Sorted | 500 | 0.693 | 3524 | 0 | Yes  
31 Quick Sort | Nearly Sorted | 500 | 1.855 | 5257 | 2850 | Yes  
32 Merge Sort | Random | 1000 | 1.428 | 8680 | 0 | Yes  
33 Quick Sort | Random | 1000 | 8.544 | 11793 | 6866 | Yes  
34 Merge Sort | Reversed | 1000 | 5.879 | 5469 | 0 | Yes  
35 Quick Sort | Reversed | 1000 | 18.934 | 15279 | 9714 | Yes  
36 Merge Sort | Nearly Sorted | 1000 | 9.196 | 7988 | 0 | Yes  
37 Quick Sort | Nearly Sorted | 1000 | 4.297 | 12119 | 7083 | Yes  
38 ======  
39  
40 --- 3. Benchmark Analysis ---  
41 Merge Sort (In-Place):  
42 - Consistent  $O(n\log n)$  performance. Time scales predictably with size.  
43 - Performance is almost identical for Random, Reversed, and Nearly Sorted data,  
44 | proving it is not vulnerable to input order. Swaps are 0 as it copies.  
45  
46 Quick Sort (Median-of-Three):  
47 - Fastest on average (especially for Random and Nearly Sorted data).  
48 - The Mo3 pivot strategy successfully prevented  $O(n^2)$  behavior on Reversed  
49 | and Nearly Sorted lists, keeping performance at  $O(n \log n)$ .  
50 - Shows the highest number of swaps, as it works in-place.  
51  
52 Conclusion:  
53 - For general-purpose sorting, QuickSort (with Mo3) is fastest.  
54 - If stability is required (preserving order of 30m, 30m), Merge Sort is the only choice.  
55 - If predictable, worst-case performance is critical, Merge Sort is safer.
```

6) Reflection: Key learnings and design trade-offs.

The project demonstrated that there is no single best algorithm or data structure. Every design choice made will have a trade-off. For example, average time for the hash table implemented in module 2 is a good example of trade-offs. Its great performance is only achieved because of computationally expensive, but rare `resize()` operation. Another example can be in module 3, ability to be fully sorted is traded for ability to insert in $O(\log n)$ time. Moreover, in module 4, merge sort guarantees a time complexity of $O(n \log n)$ at the cost of $O(n)$ space complexity.

7) Limitations & Assumptions: Known issues and simplifications.

1. Assuming all input files exist in the correct directory and contain the exact headers as outlined in the python code.
2. Module 1 cannot account for real time changes like for example a corridor being inaccessible.
3. Reinsertion assumes the older lower priority request will eventually be deleted.